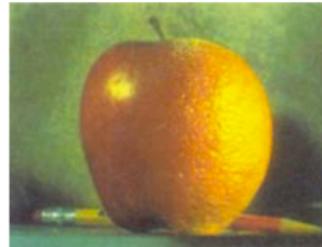


2. Image Formation



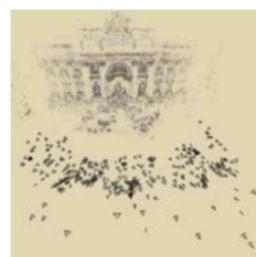
3. Image Processing



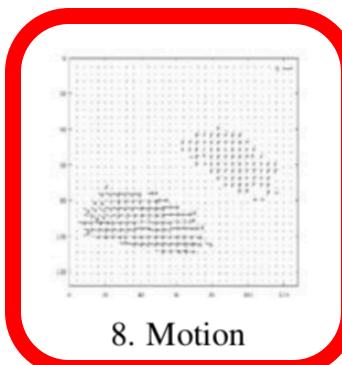
4. Features



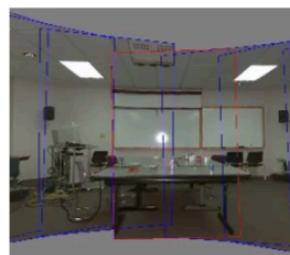
5. Segmentation



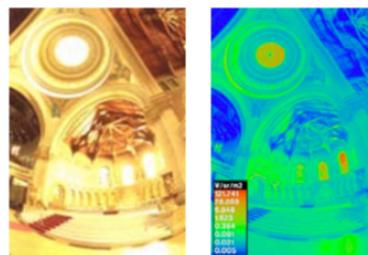
6-7. Structure from Motion



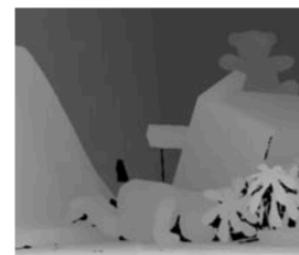
8. Motion



9. Stitching



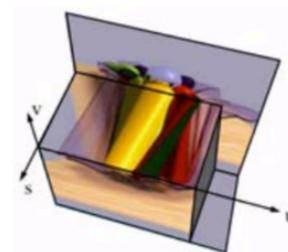
10. Computational Photography



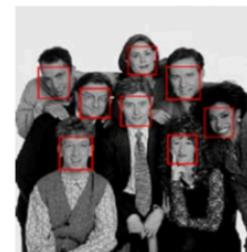
11. Stereo



12. 3D Shape



13. Image-based Rendering



14. Recognition

Credits

- Images and formulas from Szeliski
- Second half from CVPR talk by ZhaoYang Lv

Taking a Deeper Look at the Inverse Compositional Algorithm, ZhaoYang Lv , Frank Dellaert, James M. Rehg, Andreas Geiger, CVPR 2019

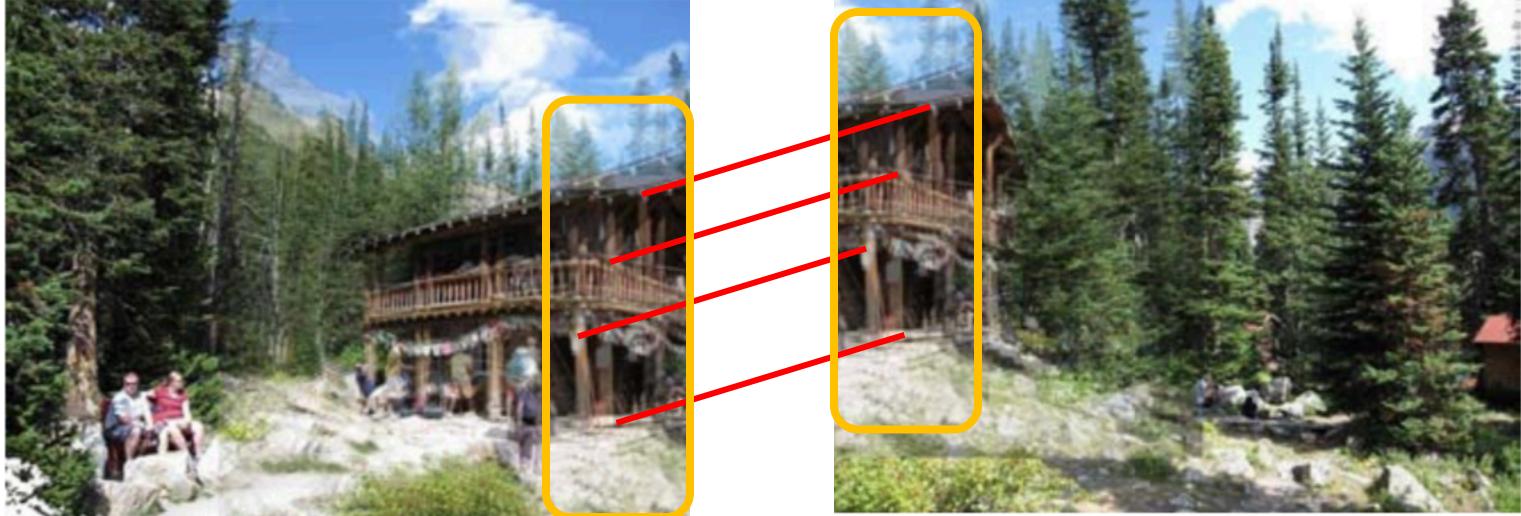
Translational Alignment



- Shift image I_1 with respect to template I_0
- Before, **feature-based** error:

$$E_{\text{LS}} = \sum_i \|r_i\|^2 = \sum_i \|f(x_i; p) - x'_i\|^2,$$

Translational Alignment



- Shift image I_1 with respect to template I_0
- Before, **feature-based** error:

$$E_{\text{LS}} = \sum_i \|r_i\|^2 = \sum_i \|f(x_i; p) - x'_i\|^2,$$

- Now, **image-based** error:

$$E_{\text{SSD}}(\mathbf{u}) = \sum_i [I_1(\mathbf{x}_i + \mathbf{u}) - I_0(\mathbf{x}_i)]^2 = \sum_i e_i^2,$$

SSD



$$E_{\text{SSD}}(\mathbf{u}) = \sum_i [I_1(\mathbf{x}_i + \mathbf{u}) - I_0(\mathbf{x}_i)]^2 = \sum_i e_i^2,$$

- Sum of Squared Differences
- Assumes: brightness constancy
- If \mathbf{u} fractional: interpolation needed
 - Bilinear (fast, good)
 - Bicubic (slower, slightly better)

Robust Error Metrics



$$E_{\text{SAD}}(\mathbf{u}) = \sum_i |I_1(\mathbf{x}_i + \mathbf{u}) - I_0(\mathbf{x}_i)| = \sum_i |e_i|.$$

- Quadratic error is unforgiving!
- Absolute error (SAD): allows for outliers
- Differentiable robust error metrics exist

Dealing with Boundary Conditions

- Should not count pixels outside
- Add two “window” functions
- Windowed SSD metric:

$$E_{\text{WSSD}}(\mathbf{u}) = \sum_i w_0(\mathbf{x}_i)w_1(\mathbf{x}_i + \mathbf{u})[I_1(\mathbf{x}_i + \mathbf{u}) - I_0(\mathbf{x}_i)]^2,$$

- Invariant to overlap: Root mean square:

$$A = \sum_i w_0(\mathbf{x}_i)w_1(\mathbf{x}_i + \mathbf{u}) \quad RMS = \sqrt{E_{\text{WSSD}}/A}$$

Violations of Brightness Constancy

- Estimate Bias and Gain

$$I_1(\mathbf{x} + \mathbf{u}) = (1 + \alpha)I_0(\mathbf{x}) + \beta,$$

$$E_{\text{BG}}(\mathbf{u}) = \sum_i [I_1(\mathbf{x}_i + \mathbf{u}) - (1 + \alpha)I_0(\mathbf{x}_i) - \beta]^2$$

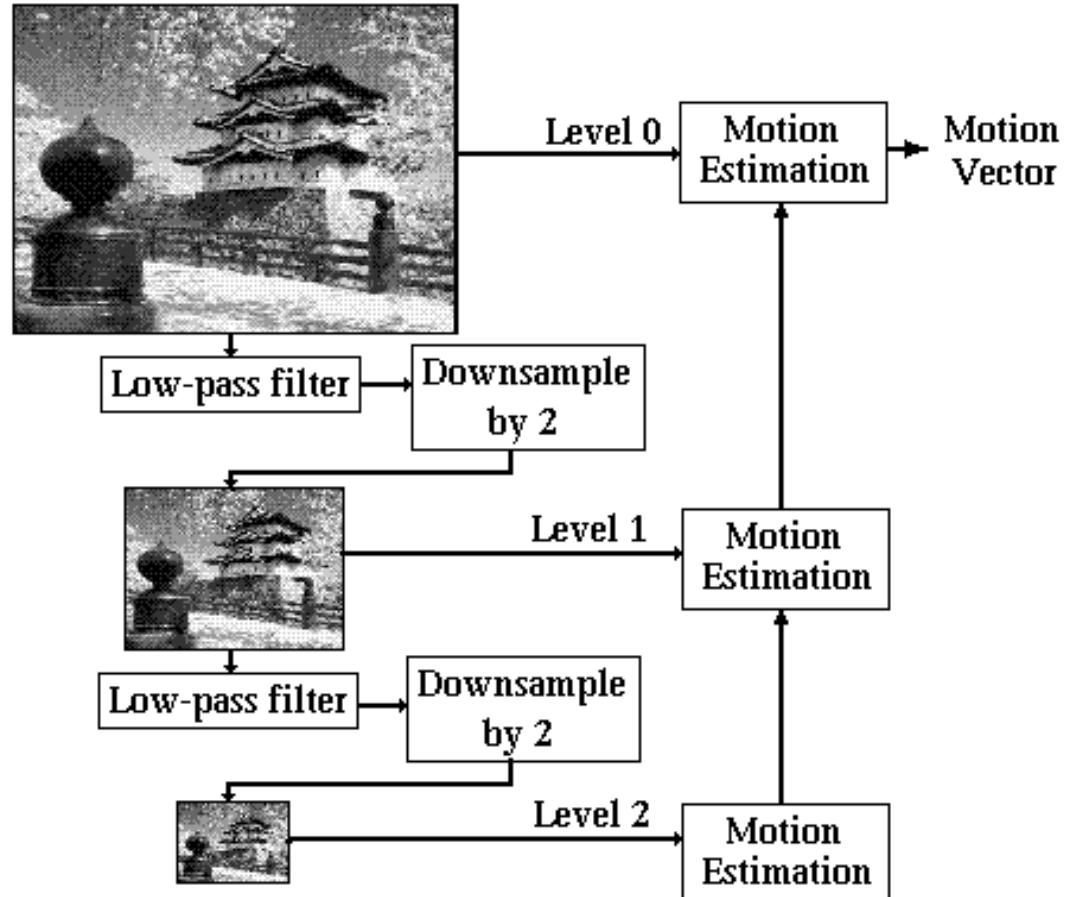
- Normalized Cross-Correlation

$$E_{\text{CC}}(\mathbf{u}) = \sum_i I_0(\mathbf{x}_i)I_1(\mathbf{x}_i + \mathbf{u}).$$

$$E_{\text{NCC}}(\mathbf{u}) = \frac{\sum_i [I_0(\mathbf{x}_i) - \bar{I}_0] [I_1(\mathbf{x}_i + \mathbf{u}) - \bar{I}_1]}{\sqrt{\sum_i [I_0(\mathbf{x}_i) - \bar{I}_0]^2} \sqrt{\sum_i [I_1(\mathbf{x}_i + \mathbf{u}) - \bar{I}_1]^2}}$$

Hierarchical Motion Estimation

- Build an image pyramid:
 - Low-pass
 - Decimate
- Recursively estimate motion:
 - Estimate motion at highest level
 - Use result as initial estimate at lower level



Sub-pixel Refinement

- Taylor expansion:

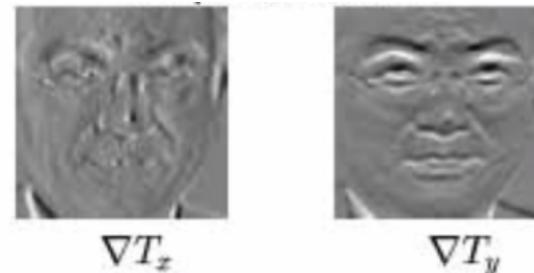
$$E_{\text{LK-SSD}}(\mathbf{u} + \Delta\mathbf{u}) = \sum_i [I_1(\mathbf{x}_i + \mathbf{u} + \Delta\mathbf{u}) - I_0(\mathbf{x}_i)]^2 \quad (8.33)$$

$$\approx \sum_i [I_1(\mathbf{x}_i + \mathbf{u}) + \mathbf{J}_1(\mathbf{x}_i + \mathbf{u})\Delta\mathbf{u} - I_0(\mathbf{x}_i)]^2 \quad (8.34)$$

$$= \sum_i [\mathbf{J}_1(\mathbf{x}_i + \mathbf{u})\Delta\mathbf{u} + e_i]^2, \quad (8.35)$$

where

$$\mathbf{J}_1(\mathbf{x}_i + \mathbf{u}) = \nabla I_1(\mathbf{x}_i + \mathbf{u}) = \left(\frac{\partial I_1}{\partial x}, \frac{\partial I_1}{\partial y} \right)(\mathbf{x}_i + \mathbf{u}) \quad (8.36)$$

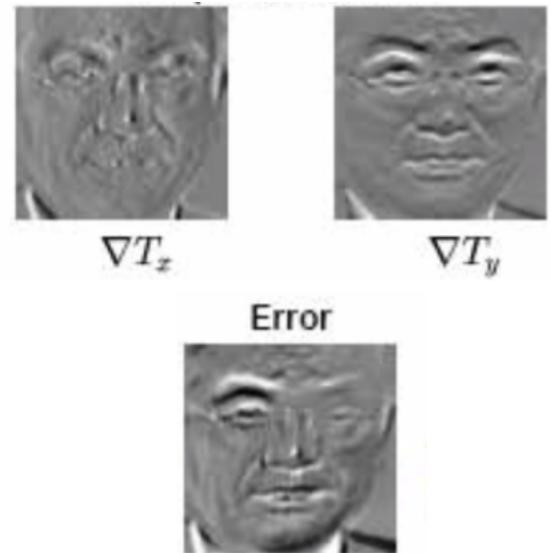


Solve using Normal Equations

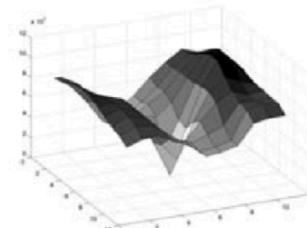
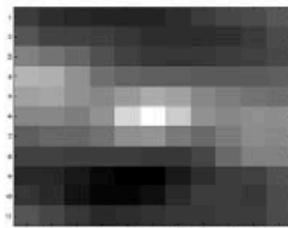
$$\mathbf{A} \Delta \mathbf{u} = \mathbf{b}$$

$$\mathbf{A} = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \text{ and } \mathbf{b} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$
$$= \sum_i \mathbf{J}_1^T(\mathbf{x}_i + \mathbf{u}) \mathbf{J}_1(\mathbf{x}_i + \mathbf{u})$$

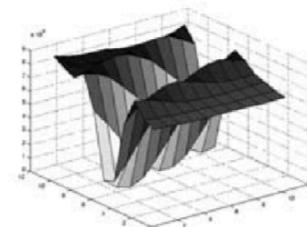
- RHS b is just dot product of gradient images with error ->
- A is Hessian or "information matrix", same as Harris uses!
- Remember: feature-based translation: just mean of flow vectors !



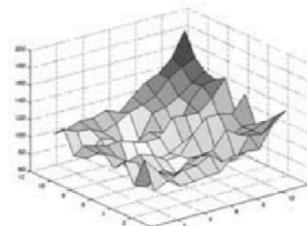
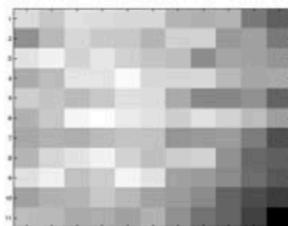
Aperture Problems and Harris



(a)



(b)



(c)

Parametric Motion

Template T



$$h(X, T) \rightarrow$$

Image I

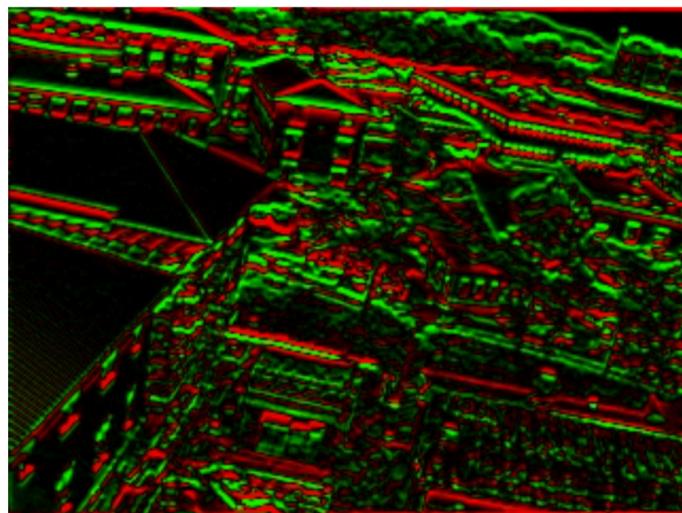
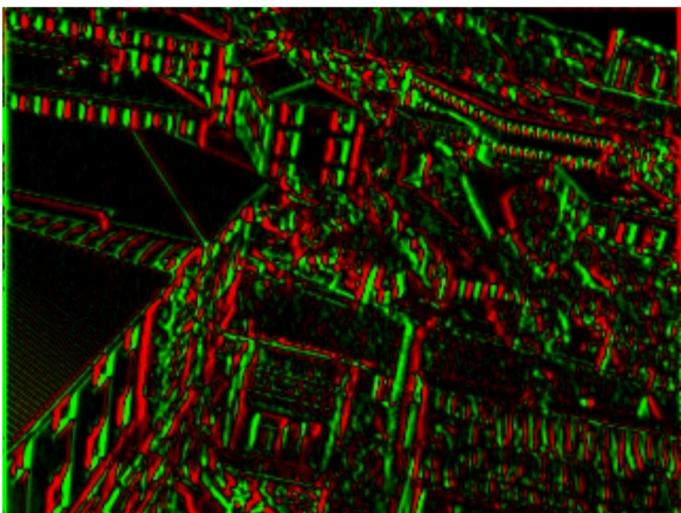
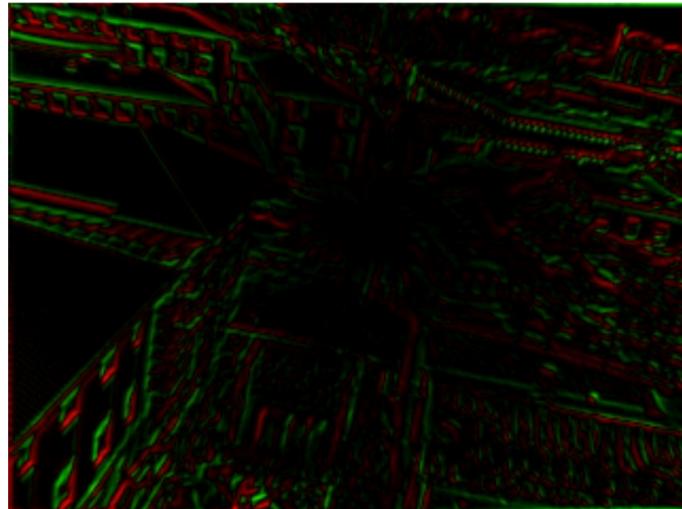


warped T = h(X, T)

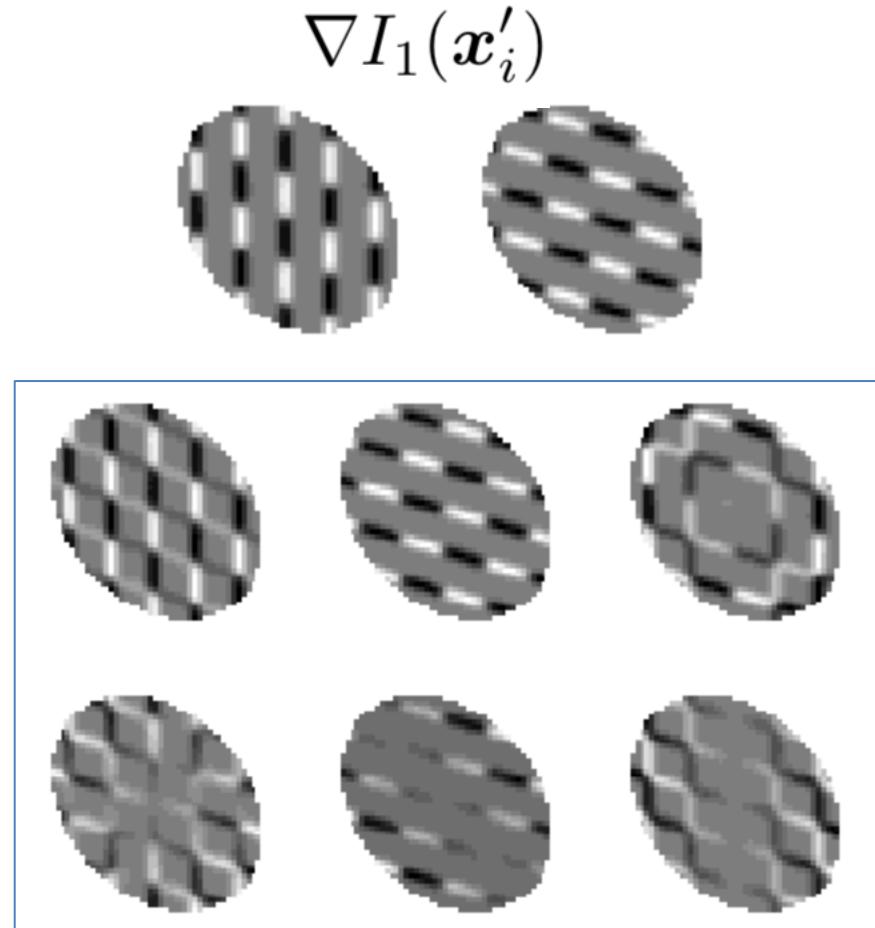
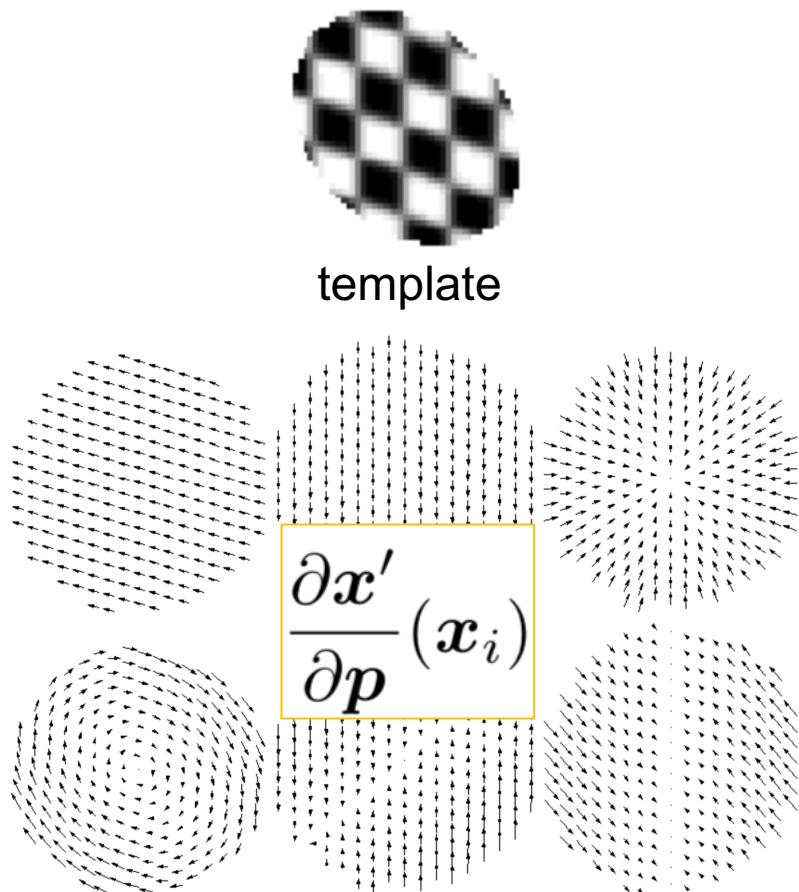
$$X = [a \ b \ c \ d \ e \ f \ g \ h]'$$

- E.g., image-based homography estimation

"Jacobian Images"

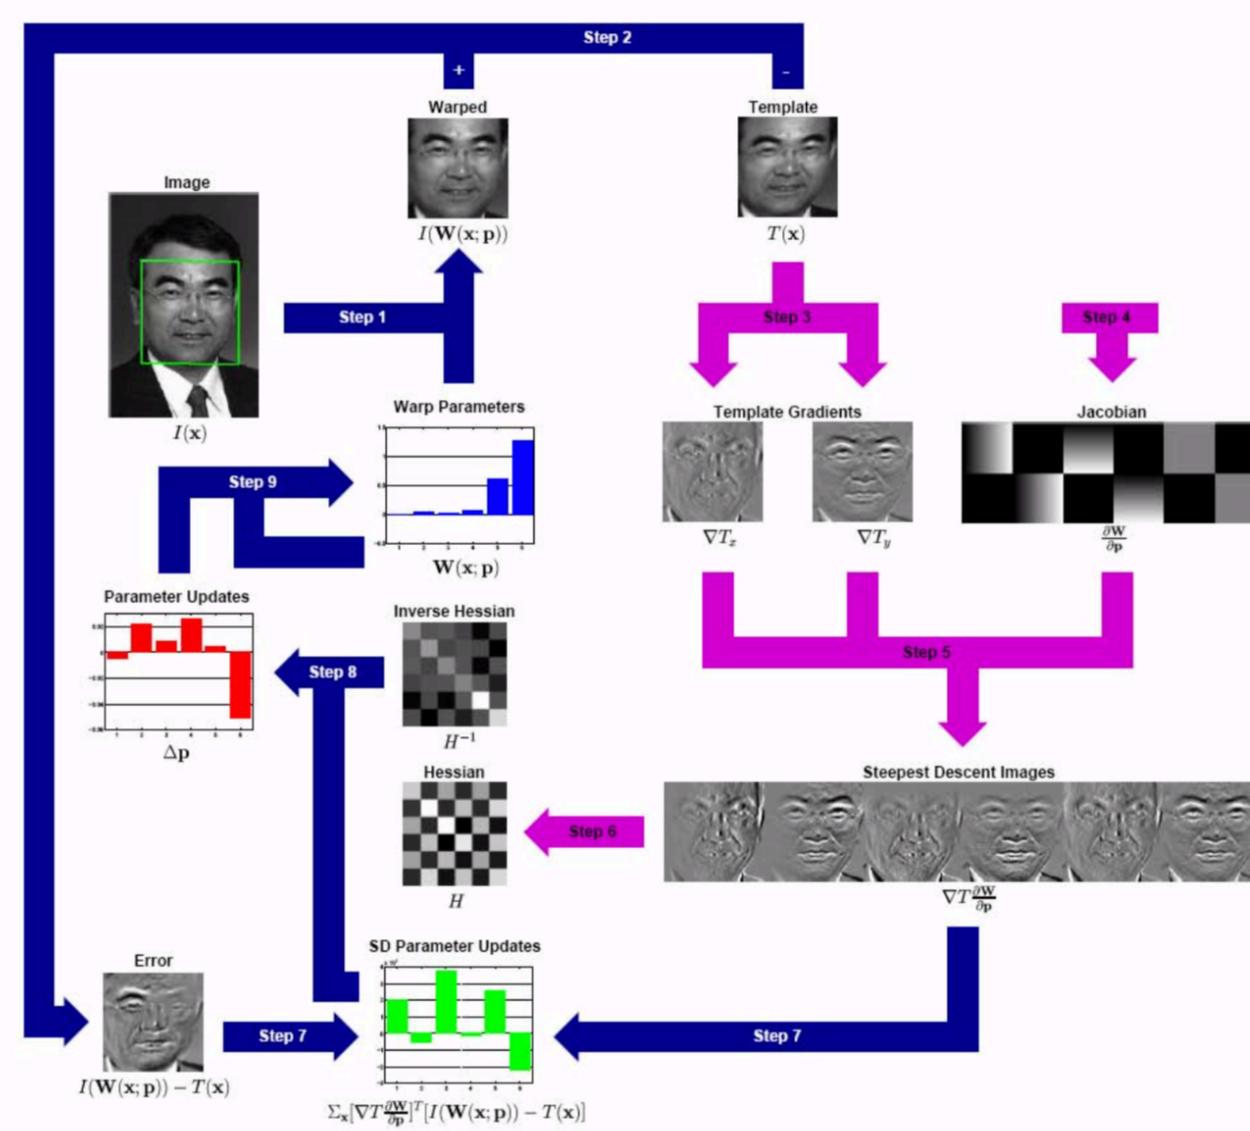


Computing Jacobian Images

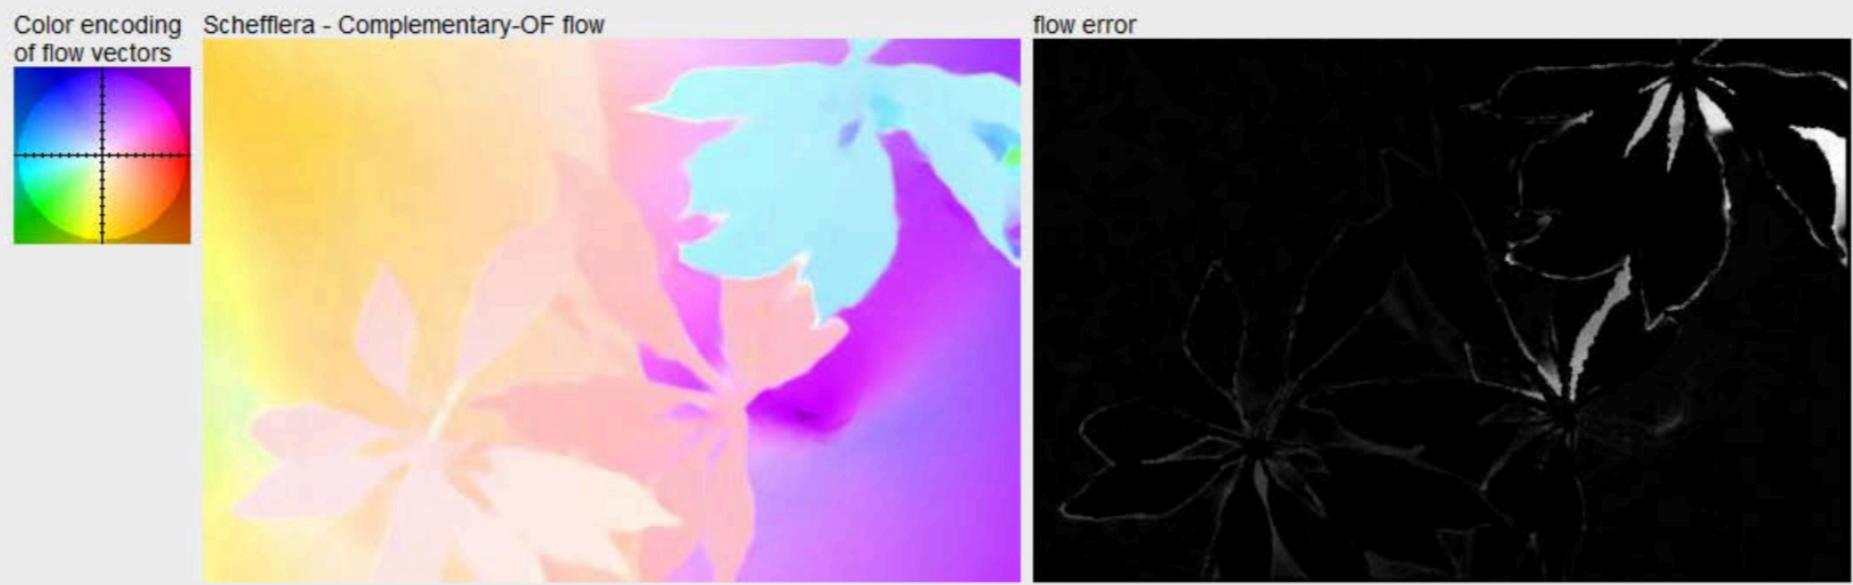


$$\mathbf{J}_1(\mathbf{x}'_i) = \frac{\partial I_1}{\partial \mathbf{p}} = \nabla I_1(\mathbf{x}'_i) \frac{\partial \mathbf{x}'}{\partial \mathbf{p}}(\mathbf{x}_i), \quad (8.52)$$

Inverse Compositional Approach

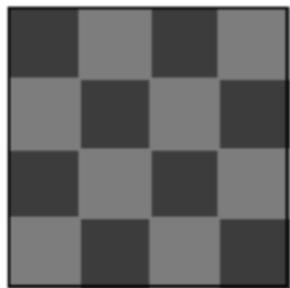


Optical Flow

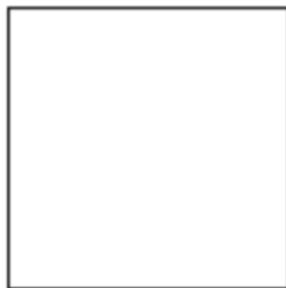


- Fully non-parametric model of motion
- N pixels $\rightarrow N$ flow vectors $\rightarrow 2N$ parameters
- Need some smoothness assumptions!
- Hard to deal with occlusion

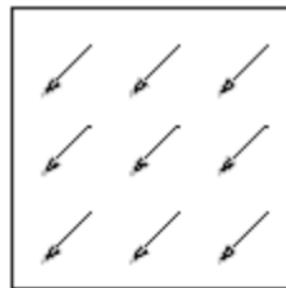
Layered Motion



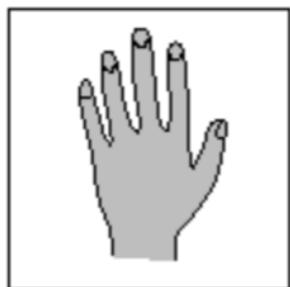
Intensity map



Alpha map



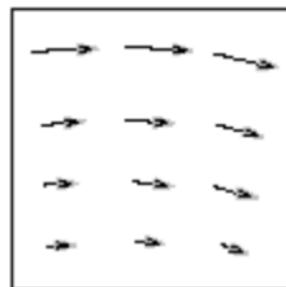
Velocity map



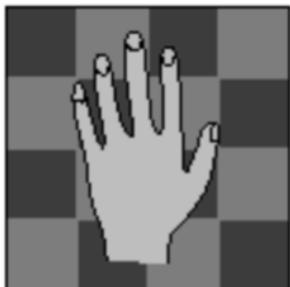
Intensity map



Alpha map



Velocity map



Frame 1



Frame 2



Frame 3

- One type of assumption to “regularize” optical flow
- Estimate FG and BG layers

Layered Motion Results



(a)



(b)



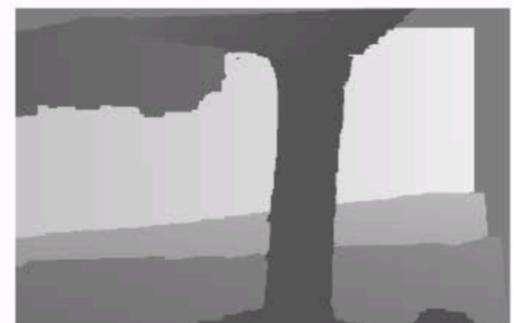
(c)



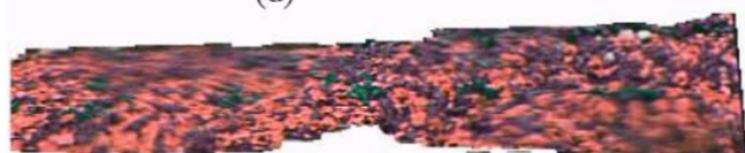
(d)



(e)



(f)



(g)



(h)

Baker, Szeliski, and Anandan 1998

The Inverse Compositional Algorithm

[S. Baker and I. Matthews, 04]

$$\rightarrow \mathbf{r}_k(\mathbf{0}) = \mathbf{I}(\boldsymbol{\xi}_k) - \mathbf{T}(\mathbf{0})$$



$$\Delta\boldsymbol{\xi} = (\mathbf{J}^T \mathbf{W} \mathbf{J} + \lambda \text{diag}(\mathbf{J}^T \mathbf{W} \mathbf{J}))^{-1} \mathbf{J}^T \mathbf{W} \mathbf{r}_k(\mathbf{0})$$

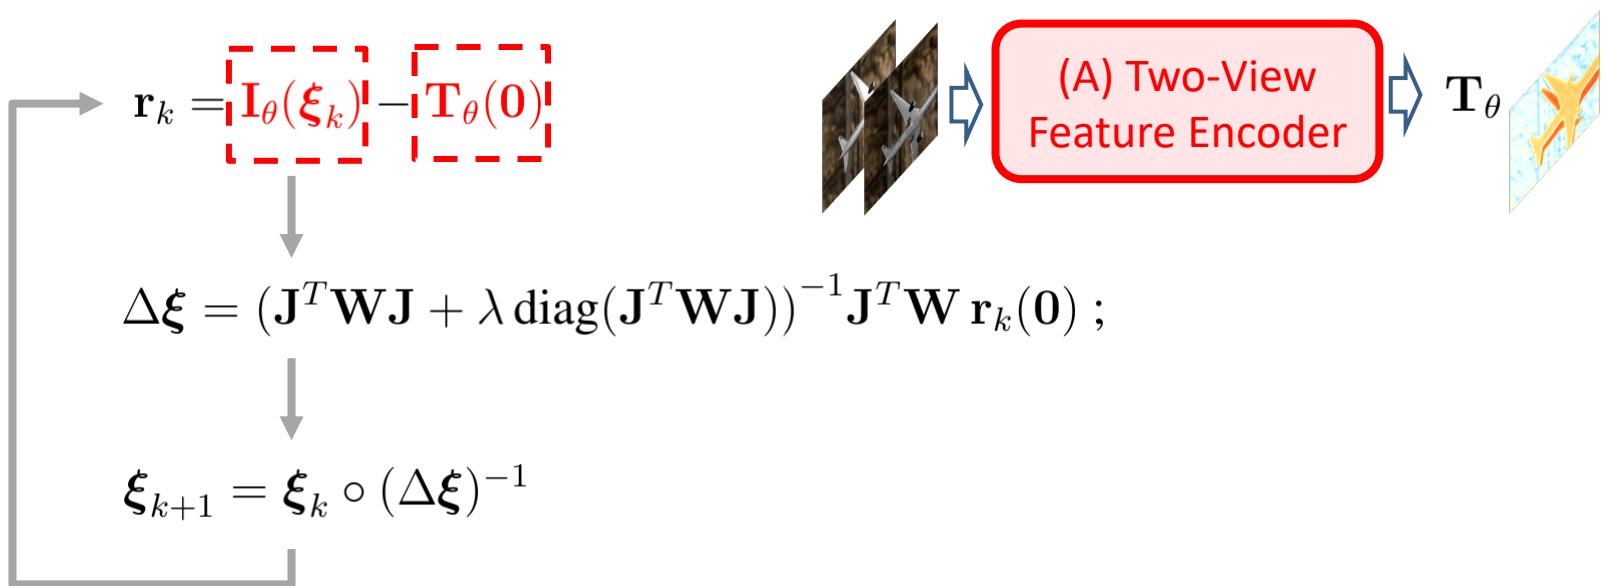


$$\boldsymbol{\xi}_{k+1} = \boldsymbol{\xi}_k \circ (\Delta\boldsymbol{\xi})^{-1}$$

We propose to take a **deeper** look at
the Inverse Compositional algorithm
from a learning perspective.

Take a Deeper Look at the Inverse Compositional algorithm

Contribution (A): Two-view Feature Encoder



Take a Deeper Look at the Inverse Compositional algorithm

Contribution (B): Convolutional M-estimator

$$\rightarrow \mathbf{r}_k = \mathbf{I}_\theta(\boldsymbol{\xi}_k) - \mathbf{T}_\theta(\mathbf{0})$$

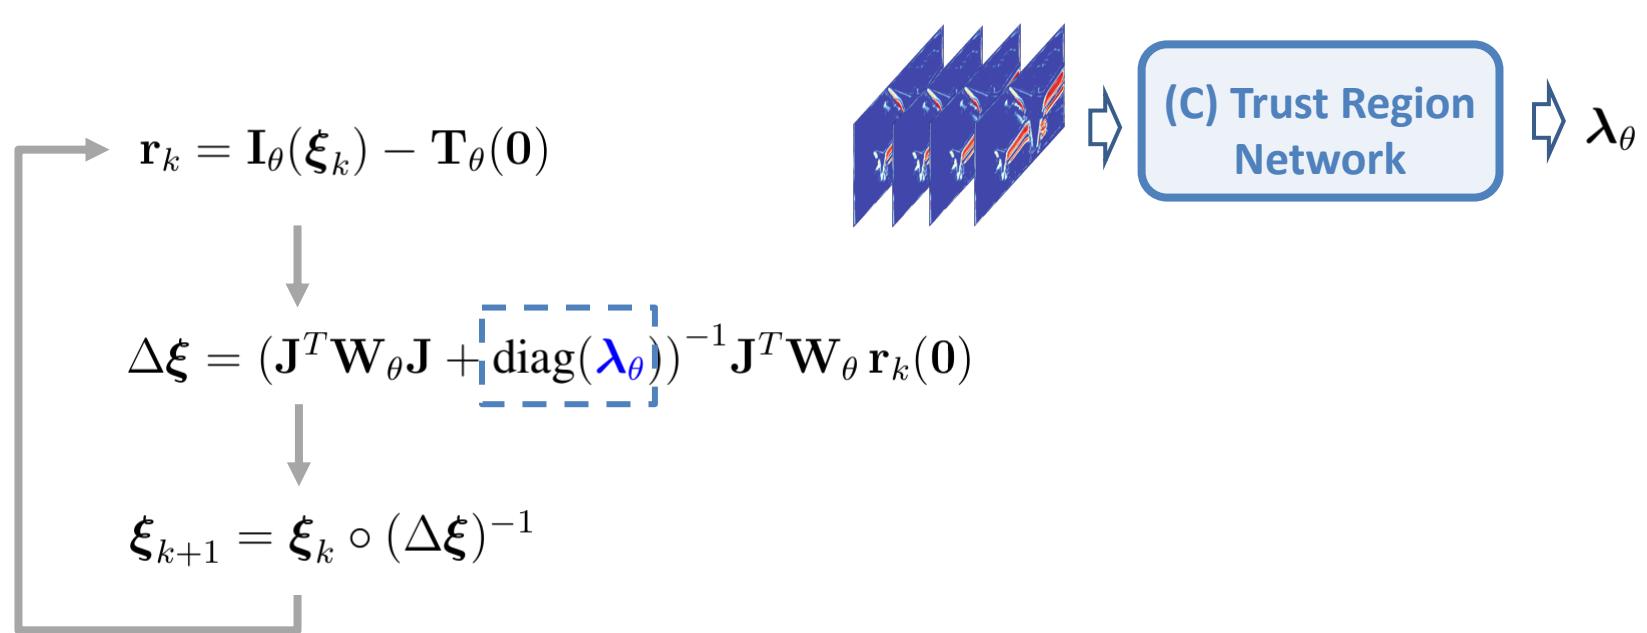
$$\Delta \boldsymbol{\xi} = (\mathbf{J}^T \boxed{\mathbf{W}_\theta} \mathbf{J} + \text{diag}(\mathbf{J}^T \boxed{\mathbf{W}_\theta} \mathbf{J})^{-1} \mathbf{J}^T \boxed{\mathbf{W}_\theta} \mathbf{r}_k(\mathbf{0}))$$

$$\boldsymbol{\xi}_{k+1} = \boldsymbol{\xi}_k \circ (\Delta \boldsymbol{\xi})^{-1}$$

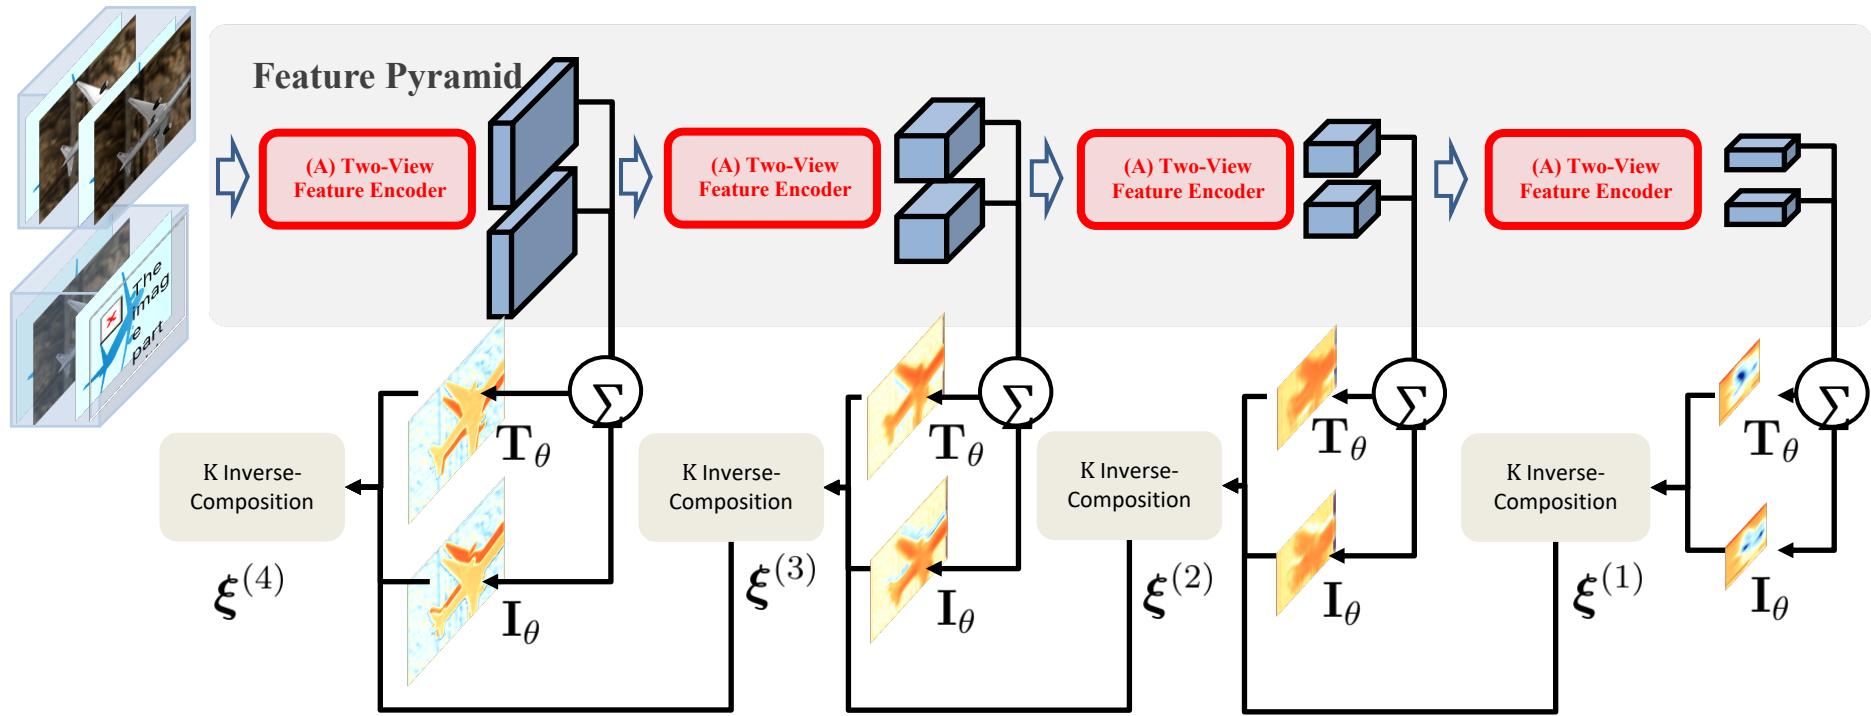


Take a Deeper Look at the Inverse Compositional algorithm

Contribution (C): Trust Region Network



Coarse-to-Fine Inverse Compositional Algorithm



Conclusion

We have taken a deeper look at the inverse compositional algorithm by reformulating it with

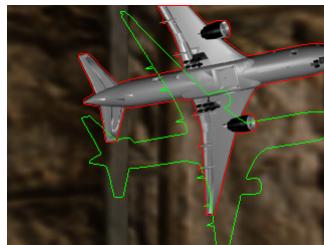
- (A) Two-view Feature Encoder
- (B) Convolutional M-estimator
- (C) Trust Region Network

The proposed solution is **learnable, accurate, small, and fast** in inference.

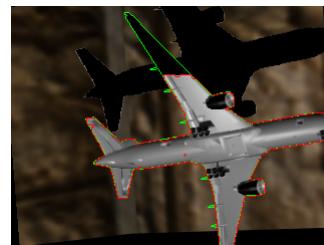
Visualization of Iterative 3D Rigid Motion Alignment



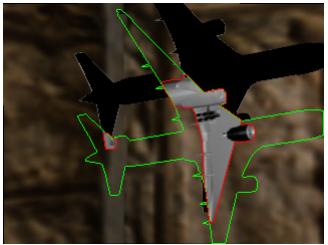
T



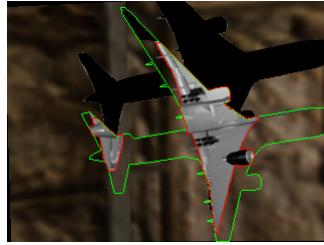
I



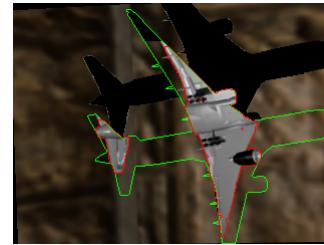
$I(\xi^{GT})$



DeepLK
[Wang et al. ICRA, 2018]



Ours (A)



Ours (A)+(B)



Ours (A)+(B)+(C)